COSC 369 · SOFTWARE ENGINEERING I · FIRST PROJECT

Professor Teaching and Office-Hours Schedule Generator

Design and implementation report

Student	[Your name]
Student ID	[Your ID]
Instructor	[Instructor's name]
Date	October 2026
Repository	`github.com/Rovers16/office-hours`

SUMMARY

This project automatically produces each professor's teaching schedule and office-hours schedule from two input files: a syllabus / class-schedule file and a faculty file. The scheduling program is written in C and organised into five modules that are archived into a static library, `libofficehours.a`, and linked with a small `main` program. A Makefile automates the whole process: compiling, building the library, linking, preparing the inputs, running the program, and generating a Word document. Python is used at the edges of the system. On the input side, Claude extracts the data from syllabi in any layout. On the output side, Python lays out the Word document and connects to a Prolog knowledge base that answers questions such as *"Who teaches COSC 369?"* and *"Which office hours are on Wednesday?"*. With the sample data, the generated office hours match the sample tables in the assignment exactly.

1. Requirements
2. System overview
3. The two input files
4. Reading syllabi in any format
5. C program design
6. Office-hour scheduling algorithm
7. Static library
8. Makefile
9. Prolog knowledge base
10. Output
11. Testing
12. Compiling and running
13. Limitations and future work

1. Requirements

The assignment asks for a C system that reads exactly two input files and generates a professor teaching schedule and an office-hours schedule. The output should look like the sample tables in the specification, without hard-coding them. The solution must:

- be organised into modules, with a **static library** that the main program links against;
- provide a **Makefile** that compiles the code, creates the library, links it, and generates the final output, optionally using Python to produce Word or Excel documents;
- use Python to read the syllabi and generate **Prolog facts**, with Prolog storing the knowledge and answering questions;
- explain static libraries, the Makefile, and where the input files are stored and why.

An additional goal was that the system should read a syllabus *however it is formatted*. Real syllabi mix tables, paragraphs and abbreviations such as "MWF 11:00–11:50am", so a parser written for one layout would fail on the next. Section 4 explains how this is handled.

2. System overview

The system is a pipeline of four stages. Each stage reads files and writes files, so the stages can be built, tested and replaced independently, and `make` can tell which stages are out of date.

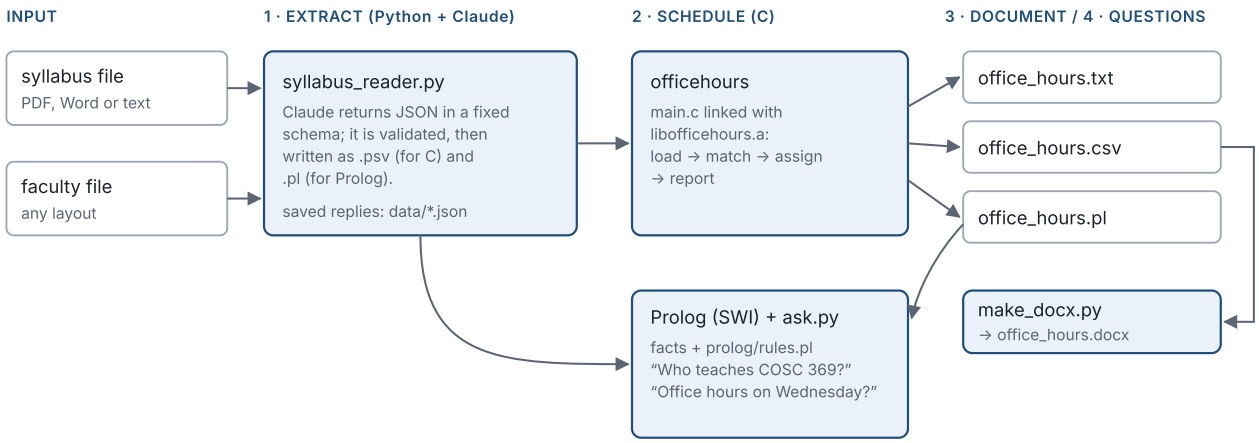


Figure 1. Data flow. Rounded blue boxes are programs; white boxes are files. The schedule.pl and faculty.pl facts from stage 1, and the office_hours.pl facts from stage 2, are all loaded by Prolog.

The responsibilities are split deliberately. **All scheduling decisions are made in C**, as the assignment requires. Python is used only where it is clearly the better tool: talking to an AI service and parsing its JSON, producing a `.docx` file, and bridging to Prolog. Every stage writes plain text files. This makes each one easy to inspect by hand, which helped a great deal during debugging.

3. The two input files

The C program, `officehours`, takes exactly two input files. It refuses to run with any other number:

```
build/officehours [options] SCHEDULE_FILE FACULTY_FILE
```

Both files are pipe-separated text, one record per line. This format is trivial to parse in C (`fgets` plus a field splitter) and easy to read and edit by hand. Lines starting with `#` are comments.

```
# schedule file
COURSE|code|title|section|lecturer|days|start|end|room
COURSE|COSC 369|Software Engineering I||Dr. Mensah|Tuesday,Thursday|09:30|10:45|CS 108

# faculty file
PROF|name|department|email|office|course,course,...
PROF|Dr. Linda Okafor|Computer Science|lokafor@university.edu|CS 214|COSC 110,COSC 241,COSC
450
OFFICE_HOUR|name|day|start|end|course
OFFICE_HOUR|Dr. Linda Okafor|Monday|13:00|15:00|COSC 241
```

Times use 24-hour `HH:MM` and are stored internally as minutes after midnight, so comparing and adding times is plain integer arithmetic. Empty fields are allowed: a course with no lecturer or no meeting time, for example an online course, is valid. `strtok()` was *not* used to split fields, because it merges consecutive separators and would silently shift every field after an empty one. A small splitter in `loader.c` keeps empty fields instead. Every line is validated, and errors are reported with the file name and line number, for example `bad.psv:1: bad time '25:00' (expected HH:MM)`.

The faculty file is also where a lecturer can supply their own office-hours table. The specification says the program should process "office-hours tables provided by the lecturer". When a professor has `OFFICE_HOUR` lines, those hours are used. Otherwise the program generates office hours itself (Section 6).

3.1 Local files or the Internet?

The input files are **stored locally** in the project: the original human-written files in `input/` and the structured data extracted from them in `data/`. The program opens them with ordinary `fopen()` calls on paths given on the command line, and the Makefile passes relative paths inside the project.

The reasons for this choice are:

- **Reliability on the lab computer.** `make` works with no Internet connection and no account or API key. A department web page could be down, moved, or reformatted on the day of the demonstration.
- **Reproducibility.** The same inputs always give the same output, which is what makes the automated test in Section 11 possible.
- **Simplicity and safety in C.** Downloading would need an HTTP library such as libcurl, extra error handling, and code to parse untrusted remote content inside the C program.

The Internet is used in exactly one optional place: turning a *new* syllabus into structured data with Claude (Section 4). That happens once per syllabus, and the result is saved locally. With Ollama mode, even that step

runs without the Internet once the model has been downloaded. A department that publishes its schedule online would simply download the file into `input/` and re-run that step.

4. Reading syllabi in any format

Syllabi do not share a format. The sample schedule in the assignment is a table, while other courses describe their meetings in sentences ("COSC 450 Operating Systems meets TTh 2:00-3:15 PM in CS 108"). Rather than writing fragile parsing rules, `syllabus_reader.py` gives the file to Claude, a large language model, and asks for the data back in one **fixed JSON structure**. The structure is defined once in Python with Pydantic models:

```
class Course(BaseModel):
    code: str                   # "COSC 369"
    title: str
    section: Optional[str]
    lecturer: Optional[str]
    days: List[Day]             # "Monday" ... "Sunday" only
    start_time: Optional[str]   # "09:30" (24-hour)
    end_time: Optional[str]
    room: Optional[str]
    notes: Optional[str]
```

Claude is instructed to expand abbreviations (`MWF` → Monday, Wednesday, Friday), convert times to 24-hour format, and use `null` rather than invent anything the file does not state. Whatever comes back is **validated against the schema** before it is used. A missing field or an invalid day name is reported as an error instead of flowing silently into the C program. The validated data is then written in three formats: `.psv` for C, `.pl` facts for Prolog, and `.json`.

There are three ways to run this step. All of them produce the same kind of files, so the rest of the system does not know or care which one was used:

Mode	How it works	Cost / requirements
Chat mode (default)	`syllabus_reader.py syllabus --prompt` prints the instructions and the JSON schema. The user pastes them into the claude.ai chat with the file attached, and saves the reply in `data/`. `--from-json` validates the saved reply and converts it; extra text or code fences around the JSON are tolerated.	Free claude.ai account; one manual step per syllabus.
API mode (`make MODE=api`)	The file is sent to the Claude API: PDF files directly, Word files as extracted text including tables. Structured outputs force the reply to match the schema.	An Anthropic API key; billed per request.
Ollama mode (`make MODE=ollama`)	An open model runs on the user's own computer through Ollama. Local models read text, so PDF and Word files are converted to text first (`pypdf`, `python-docx`). The reply is constrained to the same JSON schema and validated in the same way. The model can be chosen with `--model`.	Free; no account or Internet after the model download. An 8-billion-parameter model needs at least 8 GB of RAM (16 GB is comfortable).

Chat mode exists because the free claude.ai plan has no programming interface, and automating the website would break its terms of use and break whenever the site changes. Splitting the work into "Claude reads the file"

and “the program processes the result” keeps the system free to run while still handling any syllabus layout. Ollama mode removes the manual step and the need for an account, at the price of accuracy: an 8-billion-parameter local model is more likely than Claude to misread an unusual layout, and it cannot read scanned PDFs, which have no text layer. Because every reply passes the same validation, a model that gets the structure wrong produces an error naming the missing fields, never bad input for the C program.

5. C program design

The C code is split into modules by responsibility. Each module has a header in `include/` that declares only what other modules need. Functions used only inside a module are `static`.

Module	Responsibility	Main functions
`model.c / .h`	Shared data structures (`Course`, `Professor`, `OfficeHour`, `Schedule`, `Faculty`); growable arrays; comparing course codes and person names.	`schedule_add`, `faculty_add`, `same_course_code`, `same_person`
`timeutil.c / .h`	Converting between text and numbers: `"09:30"` ↔ 570 ↔ `"9:30 AM"`, day names, day lists such as `"Monday, Wednesday & Friday"`.	`parse_time`, `format_time`, `parse_day_list`, `format_day_list`
`loader.c / .h`	Reading and validating the two input files.	`load_schedule`, `load_faculty`
`scheduler.c / .h`	Matching professors to course sections; generating office hours or checking the ones lecturers provided.	`match_courses`, `assign_office_hours`, `default_policy`
`report.c / .h`	Writing the results as a text report, a CSV file and Prolog facts.	`write_text_report`, `write_csv_report`, `write_prolog_facts`
`main.c`	Command-line options, opening files, calling the library in order, exit status.	`main`

The dependencies between modules point one way only. `model` depends on nothing; `timeutil` uses `model`; `loader`, `scheduler` and `report` use both; `main` uses all of them. There are no cycles, so each module can be understood and tested on its own.

5.1 Data representation

- **Times** are integers (minutes after midnight). `NO_TIME` (-1) marks an online course or a time that has not been announced.
- **Meeting days** are a 7-bit mask (bit 0 = Monday), so “does this course meet on Wednesday?” is a single `&`.
- **The lists of courses and professors** grow with `realloc` (doubling), so there is no fixed limit on the size of a department. Per-professor lists have fixed maximums (16 sections, 32 office-hour blocks). Exceeding them is reported, never overflowed. All string copies go through `copy_field`, which always terminates the string and never writes past the buffer.

5.2 Matching professors to courses

The syllabus and the faculty file are written separately, so names rarely match exactly. The syllabus may say “Dr. Mensah” while the faculty file says “Dr. Kwame Mensah”. `same_person` splits both names into lower-case words, ignores titles such as Dr. and Prof., and accepts the match when every word of the shorter name appears in the longer one. A section belongs to a professor if the syllabus names them as its lecturer. If the syllabus names nobody, the professor's course list in the faculty file decides. Course codes are compared ignoring case, spaces and punctuation, so `COSC 369`, `cosc369` and `COSC-369` are the same course. A section nobody teaches is reported as a warning rather than silently dropped.

6. Office-hour scheduling algorithm

The specification's sample shows the pattern to follow. Office hours start when a class ends (10:45 AM after a 9:30–10:45 class), and an evening class gets office hours *before* it (3:00–5:00 PM before a 5:00 PM class). It also states that instructors are not to be interrupted during class time. The algorithm turns these observations into rules.

Lecturer-supplied hours. If the faculty file contains office hours for a professor, they are kept exactly as given. Each one is checked against that professor's classes, and any overlap is printed as a warning.

Generated hours. Otherwise, every class meeting is processed in week order (day, then start time), so morning blocks are placed before afternoon ones. For each meeting:

1. Build the professor's *busy list* for that day: their classes plus the office hours placed so far.
2. **After the class:** find how long the professor is free from the moment the class ends, up to the next busy item or the end of the day (5:00 PM). Use up to 135 minutes of that gap.
3. **Before the class**, if there was no room after it: find how long they are free before the class starts, back to the previous busy item or 8:00 AM. Use up to 120 minutes of that gap.
4. Gaps shorter than 30 minutes are not worth scheduling. If neither gap is long enough, a warning is printed for that meeting.

Because the busy list always includes the classes, a generated block can never overlap a class or another block. All five numbers (135, 120, 30, 8:00, 17:00) live in a `Policy` structure and can be changed from the command line, for example `--after 120 --day-end 18:00`, without recompiling.

6.1 Worked example

Dr. Mensah teaches COSC 330 (MW 9:30–10:45), COSC 369 (TTh 9:30–10:45) and COSC 631 (W 5:00–7:30 PM).

Meeting	Decision	Block
Mon COSC 330	Free from 10:45 until 5:00 PM; take 135 min.	10:45 AM – 1:00 PM
Tue / Thu COSC 369	Same reasoning.	10:45 AM – 1:00 PM
Wed COSC 330	Same reasoning.	10:45 AM – 1:00 PM
Wed COSC 631	Ends at 7:30 PM, after the 5:00 PM limit, so there is no room after. Before: free since 1:00 PM (the end of the previous block); take 120 min ending at 5:00 PM.	3:00 – 5:00 PM

This reproduces the sample office-hours table in the assignment exactly, including the Wednesday "& 3:00 – 5:00 PM (COSC 631)" line and the empty Friday. The table was not copied into the program; it follows from the rules.

7. Static library

A **static library** is an archive of compiled object files (`.o`) with an index of the functions they define. On Unix it is named `lib`*`name`*`.a`. When a program is linked against it, the linker copies the object code the program actually uses into the executable. The finished program therefore does not need the library file at run time.

Creating it. Each library module is compiled on its own into an object file, without linking (`-c`). The objects are then archived with `ar`:

```
gcc -Iinclude -std=c99 -Wall -Wextra -pedantic -O2 -c src/model.c -o build/model.o
...                                                  (timeutil, loader, scheduler, report)
ar rcs build/libofficehours.a build/model.o build/timeutil.o build/loader.o \
                              build/scheduler.o build/report.o
```

The flags mean: `r` inserts or replaces members, `c` creates the archive without a warning, and `s` writes the symbol index the linker uses to find functions. `ar t build/libofficehours.a` lists the five members, and `nm` shows the public functions they provide (`load_schedule` , `assign_office_hours` , `write_text_report` , …).

Linking it.

```
gcc build/main.o -Lbuild -lofficehours -o build/officehours
```

`-L build` adds a directory to the linker's search path, and `-lofficehours` means "look for `libofficehours.a`". Order matters: the library is listed *after* `main.o`. The linker only takes from an archive the members that resolve symbols it is still missing, so it must already have seen `main.o`'s calls by the time it reaches the library.

Static versus shared. A shared library (`.so`) is not copied into the program; it is loaded when the program starts. Shared libraries save disk space and memory when many programs use them, and they can be updated without relinking. However, the program then needs the `.so` to be present at run time. For a self-contained program that must run on a lab computer, a static library is simpler and more robust. Building the library separately also makes the scheduling code reusable. A different front end, such as a web service or a unit-test driver, could link the same `libofficehours.a` without touching `main.c`.

8. Makefile

A Makefile is a list of **rules**. Each rule names a *target* (a file to produce), the *prerequisites* it depends on, and the *commands* that rebuild it:

```
$(LIB): $(LIB_OBJS)
        $(AR) $(ARFLAGS) $@ $^          # $@ = the target, $^ = all prerequisites
```

`make` builds the requested target by walking this dependency graph. A target is rebuilt only if it is missing or *older* than one of its prerequisites. The result is that after a change, only the affected steps run. For example, after editing `timeutil.h`, `make` recompiles the five source files that include it (`model.c` does not), re-archives the library, relinks, and regenerates the reports. Running `make` again immediately prints "Nothing to be done". To make this work, every object file lists the headers it includes:

```
$(BUILD_DIR)/scheduler.o: $(SRC_DIR)/scheduler.c include/scheduler.h include/model.h
include/timeutil.h
$(BUILD_DIR)/%.o: $(SRC_DIR)/%.c | $(BUILD_DIR)
        $(CC) $(CPPFLAGS) $(CFLAGS) -c $< -o $@
```

The second rule is a *pattern rule* that compiles any `.c` file into a `.o`. The `| $(BUILD_DIR)` part is an *order-only* prerequisite: the directory must exist, but its timestamp changing does not force a rebuild.

Automating the whole pipeline. The rules continue past compilation, so a single `make` carries the work from the input files to the final document:

```
data/*.json          ➔ syllabus_reader.py ➔ build/schedule.psv, build/faculty.psv
src/*.c              ➔ gcc -c ➔ *.o ➔ ar ➔ libofficehours.a ➔ link ➔ officehours
officehours + inputs ➔ output/office_hours.txt, .csv, .pl
office_hours.csv     ➔ make_docx.py ➔ output/office_hours.docx
```

Command	What it does
`make`	Everything: library, program, reports, Word document, Prolog facts.
`make lib`	Only `build/libofficehours.a`.
`make run`	Prints the report on the screen.
`make ask`	Answers the sample questions with Prolog.
`make test`	Regenerates the outputs and compares them with the expected files.
`make MODE=api`	Reads `input/` through the Claude API instead of the saved replies in `data/`.
`make clean`	Removes `build/` and `output/`.

Variables such as `CC`, `CFLAGS` and `PYTHON` are set with `?=`, so they can be overridden on the command line. For example, `make CFLAGS="... -Werror"` makes warnings fatal, which is how the code was checked to compile without any warnings. One program run writes three output files. To keep this compatible with older versions of GNU make, the text report is the real target and the CSV and Prolog files depend on it, rather than using the newer "grouped targets" syntax.

9. Prolog knowledge base

Prolog stores knowledge as **facts** and **rules**, and answers questions by searching for values that make a query true. In this project the facts are generated, never typed in by hand:

File	Written by	Facts
`build/schedule.pl`	`syllabus_reader.py`	`course(Code, Title, Section, Lecturer, Room)`, `meets(Code, Section, Day, Start, End)`
`build/faculty.pl`	`syllabus_reader.py`	`professor(Name, Dept, Email, Office)`, `teaches(Name, Code)`, `office_hour(...)`
`output/office_hours.pl`	the C program (`-p`)	`assigned(Prof, Code, Section)`, `scheduled_office_hour(Prof, Day, Start, End, Code)`

Taking the office hours from the C program's output, rather than from the faculty file, means Prolog reasons about the *final* schedule, including generated hours. Prolog's answers and the printed report can never disagree. The hand-written rules in `prolog/rules.pl` build on these facts. Some examples:

```
% Who teaches this class?
who_teaches(Code, Prof) :- assigned(Prof, Code, _).

% Which office hours are on a given day?
office_hours_on(Day, Prof, Start, End, Code) :-
    scheduled_office_hour(Prof, Day, Start, End, Code).

% Can a student see this professor at this time?  (in office hours and not teaching)
available(Prof, Day, Time) :-
    in_office(Prof, Day, Time, _),
    \+ in_class(Prof, Day, Time).

% Does any office hour overlap one of the professor's classes?
clash(Prof, Day, Code, OStart, OEnd) :-
    scheduled_office_hour(Prof, Day, OStart, OEnd, _),
    class_time(Prof, Code, Day, CStart, CEnd),
    overlaps(OStart, OEnd, CStart, CEnd).
```

Other rules answer where a course meets, which days a professor is off campus, how many office hours they hold per week, and which courses are online or meet in the evening. Times are stored as zero-padded `'HH:MM'` atoms, so Prolog's standard ordering (`@<`) compares them correctly without any conversion.

Python ↔ Prolog. `ask.py` uses pyswip to start SWI-Prolog, consults `prolog/knowledge_base.pl` (which loads all the facts and rules), runs queries, and prints the answers. Python only builds the query and formats the result; the reasoning is done by the rules. Anything a user types is passed into the query as a quoted Prolog atom, so input like `x'), halt, ('` is treated as an unknown course name, not as code. The knowledge base can also be used directly in `swipl`.

```
$ python3 ask.py available mensah tuesday 10:00
Can I see Dr. Kwame Mensah on Tuesday at 10:00 AM?
  No - they are teaching then (instructors are not to be interrupted).

$ python3 ask.py office-hours wednesday
Which office hours are on Wednesday?
  Dr. Kwame Mensah: 10:45 AM - 1:00 PM (COSC 330)
  Dr. Linda Okafor: 1:00 PM - 3:00 PM (COSC 241)
  Dr. Kwame Mensah: 3:00 PM - 5:00 PM (COSC 631)
```

10. Output

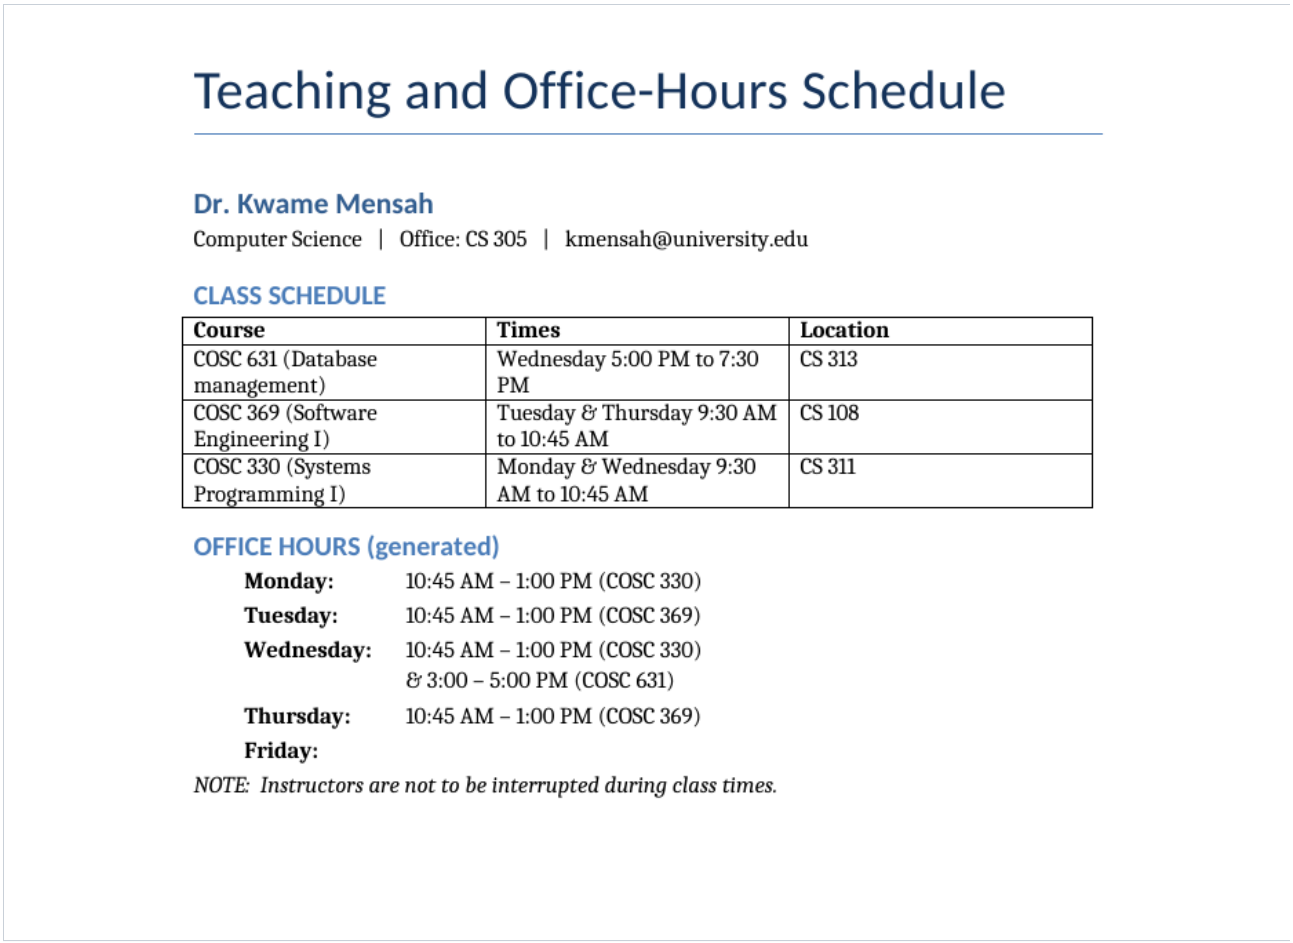
The sample data contains two professors chosen to exercise both paths through the program. Dr. Mensah has the courses from the assignment's sample and no office hours in the faculty file, so his are generated. Dr. Okafor supplies her own office hours and also teaches an online course with no meeting time. The text report follows the layout of the sample tables:

```
CLASS SCHEDULE
Course                              Times                                       Location
-------------------------------------------------------------------------------------------
COSC 631 (Database management)      Wednesday 5:00 PM to 7:30 PM                CS 313
COSC 369 (Software Engineering I)   Tuesday & Thursday 9:30 AM to 10:45 AM      CS 108
COSC 330 (Systems Programming I)    Monday & Wednesday 9:30 AM to 10:45 AM      CS 311

OFFICE HOURS (generated)
    Monday:    10:45 AM - 1:00 PM (COSC 330)
    Tuesday:   10:45 AM - 1:00 PM (COSC 369)
    Wednesday: 10:45 AM - 1:00 PM (COSC 330)
               & 3:00 - 5:00 PM (COSC 631)
    Thursday:  10:45 AM - 1:00 PM (COSC 369)
    Friday:

NOTE: Instructors are not to be interrupted during class times.
```

The same results are written to a CSV file, in which every value is already formatted for display. `make_docx.py` lays that file out as a Word document with one page per professor (Figure 2). Keeping the layout in Python and the logic in C means the document's appearance can change without touching the scheduler.



Teaching and Office-Hours Schedule

Dr. Kwame Mensah

Computer Science | Office: CS 305 | kmensah@university.edu

CLASS SCHEDULE

Course	Times	Location
COSC 631 (Database management)	Wednesday 5:00 PM to 7:30 PM	CS 313
COSC 369 (Software Engineering I)	Tuesday & Thursday 9:30 AM to 10:45 AM	CS 108
COSC 330 (Systems Programming I)	Monday & Wednesday 9:30 AM to 10:45 AM	CS 311

OFFICE HOURS (generated)

Monday: 10:45 AM – 1:00 PM (COSC 330)

Tuesday: 10:45 AM – 1:00 PM (COSC 369)

Wednesday: 10:45 AM – 1:00 PM (COSC 330) & 3:00 – 5:00 PM (COSC 631)

Thursday: 10:45 AM – 1:00 PM (COSC 369)

Friday:

NOTE: Instructors are not to be interrupted during class times.

Figure 2. First page of the generated `output/office_hours.docx`.

11. Testing

- **Regression test.** `make test` regenerates the report and the Prolog answers and compares them, using `diff`, with `tests/expected_report.txt` and `tests/expected_answers.txt`. Any change in behaviour shows up as a readable difference.
- **Edge cases.** These were tested with a separate set of inputs:
  - back-to-back classes (a gap too short after the first, so its block moves before it);
  - a class starting after 5:00 PM;
  - lecturer-supplied hours that overlap a class (reported as a warning);
  - a course whose lecturer is not in the faculty file (reported);
  - a course matched only through the faculty file's course list;
  - an online course with no meeting time.
- **Invalid input.** Unknown days, impossible times (`25:00`), a class that ends before it starts, unknown record types, a missing file, the wrong number of input files, and bad option values each produce a clear message and a non-zero exit status.

- **Memory safety.** All of the above were run with the code compiled using AddressSanitizer and UndefinedBehaviorSanitizer (`-fsanitize=address,undefined`), with no errors reported. The code also compiles cleanly with `-std=c99 -Wall -Wextra -pedantic -Werror`.
- **Python side.** The reader was tested on the assignment's own `.docx` file, including its tables, and with malformed replies that must be rejected. The Ollama path was tested against a stand-in server implementing Ollama's chat interface. This checked the request (model, schema, context size, text extracted from `.txt`, `.docx` and `.pdf` files) and every failure case: Ollama not running, model not installed, a server error, a reply cut off by the context limit, and a reply that does not match the schema. The Prolog questions were tested with loosely typed names, unknown names, and quote-injection attempts.

12. Compiling and running

```
sudo apt install build-essential python3-pip swi-prolog-nox
pip install -r requirements.txt     # pydantic, python-docx, pyswip (+ anthropic for API
mode)

make            # build the library and program, generate output/ and the Word document
make test       # verify against the expected output
make ask        # sample questions answered by Prolog
build/officehours --help
```

To use a new syllabus in chat mode:

```
mkdir -p build
python3 syllabus_reader.py syllabus --prompt > build/prompt.txt
#   paste prompt.txt into claude.ai with the syllabus attached;
#   save the reply as data/syllabus_reply.json (same for the faculty file)
make
```

13. Limitations and future work

- **Name matching is heuristic.** Two professors who share a surname would both match a syllabus entry that gives only the surname. Including first names or emails in the syllabus, or a stricter matching mode, would remove the ambiguity.
- **The office-hour policy is simple.** It does not aim for a weekly total (some departments require a fixed number of hours), avoid lunch breaks, or spread hours over days without classes. Because the policy is a separate structure, a different strategy could be added without changing the loader or the reports.
- **Extraction quality depends on the source.** A badly scanned PDF can be misread. Validation catches malformed data, but not a plausible wrong value such as a misread room number. A human should glance at the extracted data in `data/` once per syllabus.
- **Possible extensions:** an Excel export alongside the Word document; detecting room double-bookings across professors; a web front end linked against the same static library.